

PaperSwarm: An Evidence-Gated Multi-Agent System for Reproducible Paper Generation

Anonymous authors
Paper under double-blind review

Abstract

This build was produced with `-skip-llm-check`. The pipeline below (ledger, gate, assembly, compilation) is exercised, but no model-generated prose is included. Ridge improves over OLS by 50.94 percent.

Placeholder

This build was produced with `-skip-llm-check`. The pipeline below (ledger, gate, assembly, compilation) is exercised, but no model-generated prose is included.

Placeholder

This build was produced with `-skip-llm-check`. The pipeline below (ledger, gate, assembly, compilation) is exercised, but no model-generated prose is included.

Placeholder

This build was produced with `-skip-llm-check`. The pipeline below (ledger, gate, assembly, compilation) is exercised, but no model-generated prose is included.

Placeholder

This build was produced with `-skip-llm-check`. The pipeline below (ledger, gate, assembly, compilation) is exercised, but no model-generated prose is included. OLS mean test MSE is 0.6198 (standard deviation 0.1329); ridge attains 0.3041 with standard deviation 0.0623 at penalty 0.1. The relative improvement is 50.94 percent, against an irreducible noise floor of 0.1225.

Placeholder

This build was produced with `-skip-llm-check`. The pipeline below (ledger, gate, assembly, compilation) is exercised, but no model-generated prose is included.

Placeholder

This build was produced with `-skip-llm-check`. The pipeline below (ledger, gate, assembly, compilation) is exercised, but no model-generated prose is included.

References